

IDNM 680: HOMEWORK 4

JAMES DELLA-GIUSTINA

TASK 1

Using the gradient descent method, develop a python code to find the minimum of the functions:

- $f(x) = x^4 - 3x^3 + 2$
- $f(x, y) = 3(x - 2)^2 + 5(y + 3)^2$

Solution 1. Gradient descent involves finding the minimum of a function f . Let the vector of partial derivatives of f with respect to each component of x be defined as:

$$\nabla f = \left[\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right]^T$$

The gradient of a vector is the direction of steepest ascent, and taking the negative of the gradient allows us to approach some minima (if a function is convex, i.e. it is uniform continuous or Lipschitz, then any minima is a global minimum). But sometimes, we cannot find exactly what value of x yields $\nabla f = 0$. Therefore, we employ an iterative approach, by first choosing a random x_0 and evaluating $\nabla f(x_0)$. With this in hand, we find:

$$x_1 = x_0 - \eta \nabla f(x_0)$$

for a small learning rate $\eta \in (0, 1)$ usually chosen ~ 0.001 (η may also be initially larger and decrease with each iteration). In general, we have;

$$x_i = x_{i-1} - \eta \nabla f(x_{i-1}) \quad (1)$$

We can stop iterating and be safe in assuming that we have reached a minimum when $\|\eta \nabla f(x_i)\|$ falls below some specified threshold.

We first use the method of line search to decrease our learning rate to converge to a minima. In this method, we begin with a learning rate η , with parameters $\alpha \in (0, 1)$ and $\beta \in (0, 0.5)$. We calculate perform the calculation in Equation (1), while also calculating $f(x_i)$ and $f(x_{i-1})$. If $f(x_i) \leq f(x_{i-1}) - \alpha \eta \|\nabla f(x_i)\|$, then we update our learning rate $\eta = \beta \cdot \eta$. The following Python code minimizes the function $f(x) = x^4 - 3x^3 + 2$.

```
import numpy as np
import matplotlib.pyplot as plt
from numpy import linalg as LA
# Function to optimize (change this to whatever function you want to minimize)
def function(x):
    return x**4-3*x**3+2

# Line search function to find optimal learning rate
def line_search(x, grad, lr, iterations, alpha, beta, function):
    for i in range(iterations):
        cost = function(x - lr*grad)
        prev_cost = function(x)
        if cost <= prev_cost - alpha*lr*LA.norm(grad):
            return lr * beta
        else:
            return lr
    return lr

# Gradient descent with line search
```

```

def gradient_descent(x0, lr, iterations, threshold, function):
    history = [x0]
    x = x0
    step_count = 0
    beta = 0.1
    alpha = 0.5
    for i in range(iterations):
        grad = 4*x**3-9*x**2 # derivative of the function
        lr = line_search(x, grad, lr, iterations, alpha, beta, function)
        x = x - lr*grad
        history.append(x)
        step_count += 1
        if lr * LA.norm(grad) < threshold: # stopping condition
            break
    print("Number of steps taken:", step_count)
    return history

# Define initial parameters
x0 = np.random.uniform(-20,20) # starting point
lr = .001
iterations = 1000 # number of iterations
threshold = 1e-12 # threshold for stopping condition

# Run gradient descent with line search
history = gradient_descent(x0, lr, iterations, threshold, function)

# Convert the history list to a numpy array
history = np.array(history)

# Plot the function and the history of gradient descent
x = np.linspace(-50,50, 1000)
y = function(x)
plt.plot(x, y, label='Function')
plt.plot(history, function(history), 'ro-', label='Gradient descent')
plt.legend()
plt.show()

```

with convergence in 11 steps and illustrated by Figure 1.

For our function of two variables $f(x, y) = 3(x - 2)^2 + 5(y + 3)^2$, we may use the exact same code but modify it to take into account that we now have two variables to update and two components of our gradient $\nabla f(x, y)$.

```

# Function to optimize
def function(x, y):
    return 3*(x-2)**2 + 5*(y+3)**2

# Line search function to find optimal learning rate
def line_search(x, y, grad_x, grad_y, lr, function, alpha, beta, max_iterations=100,
threshold=1e-12):
    for i in range(max_iterations):
        cost = function(x - lr*grad_x, y - lr*grad_y)
        prev_cost = function(x, y)
        if cost <= prev_cost - threshold*lr*LA.norm([grad_x, grad_y]):
            return lr * beta
    else:

```




FIGURE 1. Minimizing the function $f(x) = x^4 - 3x^3 + 2$ using gradient descent and updating our learning rate through the backtrack line search method.

```

        return lr
    return lr

# Gradient descent with line search
def gradient_descent(x0, y0, lr, iterations, threshold, function):
    history_x = [x0]
    history_y = [y0]
    x = x0
    y = y0
    step_count = 0
    alpha = .1
    beta = .5
    for i in range(iterations):
        grad_x = 6*(x-2)
        grad_y = 10*(y+3)
        lr = line_search(x, y, grad_x, grad_y, lr, function, alpha, beta)
        x = x - lr*grad_x
        y = y - lr*grad_y
        history_x.append(x)
        history_y.append(y)
        step_count += 1
        if lr * np.linalg.norm([grad_x, grad_y]) < threshold: # stopping condition
            break
    print("Number of steps taken:", step_count)
    return history_x, history_y

# Define initial parameters
x0 = np.random.uniform(-10,10) # starting point for x
y0 = np.random.uniform(-10,10) # starting point for y

```

```

lr = 0.1 # initial learning rate
iterations = 1000 # number of iterations
threshold = 1e-12 # threshold for stopping condition

# Run gradient descent with line search
history_x, history_y = gradient_descent(x0, y0, lr, iterations, threshold, function)

# Convert the history lists to numpy arrays
history_x = np.array(history_x)
history_y = np.array(history_y)

# Generate a grid of (x,y) points for plotting the function surface
x = np.linspace(-10, 10, 100)
y = np.linspace(-10, 10, 100)
X, Y = np.meshgrid(x, y)
Z = function(X, Y)

# Plot the function surface and the history of gradient descent
fig = plt.figure()
ax = fig.add_subplot(projection='3d')
ax.plot_surface(X, Y, Z, cmap='inferno', alpha = .5)
ax.plot(history_x, history_y, function(history_x, history_y), 'ro-', label='Gradient
descent')
ax.plot(history_x, history_y, np.zeros_like(history_x), 'ro-', label='History of
gradient descent')
ax.set_xlabel('x')
ax.set_ylabel('y')
ax.set_zlabel('f(x,y)')
plt.legend()
plt.show()

```

This converges in 44 steps and produces Figure 2.




FIGURE 2. Minimizing the function $f(x,y) = 3(x-2)^2 + 5(y+3)^2$ using gradient descent and updating our learning rate through the backtrack line search method.

TASK 2

We will have a webcam input. How could we compute a beauty score? How to learn what is beauty? Follow the link below to create the dataset and find the range of the rating (min and max).

https://scientific-python.readthedocs.io/en/latest/notebooks_rst/6_Machine_Learning/04_Exercices/02_Practical_Work/02_CNN_1_Beauty.htmlproblem-statement

Solution 2. This code was executed in GoogleColab. We first download and unzip the data set SCUT-FBP5500 dataset <https://github.com/HCIILAB/SCUT-FBP5500-Database-Release> by the Human Computer Intelligent Interaction Lab of South China University of Technology.

```
##%capture
!pip install livelossplot
!pip uninstall gdown -y && pip install gdown
![ ! -d 'datasets' ] && echo 'datasets dir not found, will create' && mkdir datasets
![ ! -d 'logs' ] && echo 'logs dir not found, will create' && mkdir logs
![ ! -d 'drive/My Drive/Colab Notebooks/models' ] && echo 'models dir not found, will create' && mkdir 'drive/My Drive/Colab Notebooks/models'
![ ! -d './datasets/SCUT-FBP5500_v2' ] && echo './datasets/SCUT-FBP5500_v2 not found, will download & unzip it' && gdown 1w0TorBfTIqbquQVd6k3h_77ypnrvfGwf -O ./datasets/ && unzip -n -q -d './datasets/' './datasets/SCUT-FBP5500_v2.1.zip' && rm -f datasets/SCUT-FBP5500_v2.1.zip
```

In order to easily access the file system in GoogleColab, path variables are defined:

```
import os

dataset_path = os.path.relpath("datasets/SCUT-FBP5500_v2")
txt_file_path = os.path.join(dataset_path, "train_test_files", "All_labels.txt")
images_path = os.path.join(dataset_path, "Images")

assert os.path.exists(dataset_path)
assert os.path.exists(txt_file_path)
assert os.path.exists(images_path)

print("dataset_path:", dataset_path)
print("txt_file_path:", txt_file_path)
print("images_path:", images_path)
```

We then use a Pandas data frame object to read in the `All_labels.txt` CSV where a space ' ' serves as a delimiter between the two columns 'filename' and 'rating'. Pandas data frame objects have many useful built in functions described by the comments in the code:

```
import pandas as pd

df = pd.read_csv('/content/datasets/SCUT-FBP5500_v2/train_test_files/All_labels.txt',
                 sep=' ',
                 names=["filename", "rating"])

# Check the 5th first row
df.head()

# Describe the dataset
df.describe()

# Rates histogram
df.rating.hist(bins=30, density=1)
```

which gives output:

```

filename rating
0 CF437.jpg 2.883333
1 AM1384.jpg 2.466667
2 AM1234.jpg 2.150000
3 AM1774.jpg 3.750000
4 CF215.jpg 3.033333

```

```

rating
count 5500.000000
mean 2.990891
std 0.688112
min 1.016667
25% 2.500000
50% 2.833333
75% 3.533333
max 4.750000

```

and produces Figure 3. The range of ratings up to two significant digits is $[1.02, 4.75]$ and the average rating is 2.9.




FIGURE 3. Distribution of ratings and frequency.

TASK 3

Use the gradient descent method to find the minimum of $f(x, y, z) = -(x-1)^2 + 3(y-2)^3 + 4(z+3)^2$. Dynamically show the step-by-step optimization process.

Solution 3. Rather than trying to start immediately coding this, let us first do perform some analysis on our function.

Claim: The function $f(x, y, z) = -(x-1)^2 + 3(y-2)^3 + 4(z+3)^2$ has domain \mathbb{R}^3 and co-domain \mathbb{R} .

Proof. To show this, we only need to examine the term $(y-2)^3$ which is $\mathcal{O}(y^3)$ (in fact the whole function is $\mathcal{O}(y^3)$). For all $y \in \mathbb{R}$, $y^3 \in \mathbb{R}$, and therefore the co-domain of $f(x, y, z)$ is \mathbb{R} . As $|y| \rightarrow \infty$, this term dominates and causes the function to become unbounded. Therefore, there exists no global minima for the function. \square

We now provide an alternative proof which will work in general to determine whether or not an arbitrary function is convex and therefore has a global minima.

Proof. The Hessian matrix of a function is given as

$$H = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

Definition 1. An $n \times n$ symmetric real matrix M is said to be positive semi-definite iff $\mathbf{x}^\top M \mathbf{x} \geq 0$ for all \mathbf{x} in \mathbb{R}^n . Formally,

$$M \text{ is positive semi-definite} \Leftrightarrow \mathbf{x}^\top M \mathbf{x} \geq 0 \text{ for all } \mathbf{x} \in \mathbb{R}^n$$

If the Hessian matrix of a function is semi-positive definite, then the function is convex. The Hessian of our function is given as:

$$\begin{bmatrix} -2 & 0 & 0 \\ 0 & 18(y-2) & 0 \\ 0 & 0 & 8 \end{bmatrix}$$

Now take the vector $\mathbf{x}^\top = e_1 = (1, 0, 0)$. Then $\mathbf{x}^\top H \mathbf{x}$

$$\begin{bmatrix} 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} -2 & 0 & 0 \\ 0 & 18(y-2) & 0 \\ 0 & 0 & 8 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = -2 \not\geq 0$$

Therefore the function is non-convex and consequentially there does not exist any global minima. However, we can now talk about finding a local minimum in terms of eigenvalues. H may not be positive definite for any vector $\mathbf{x} \in \mathbb{R}^3$, but it may have a local minimum if all eigenvalues of H are positive. Since our Hessian H is diagonal, then all eigenvalues lie on the main diagonal and so $\lambda_i = \{-2, 18(y-2), 8\}$ for $i \in [3]$. This implies that for $\forall y > 2$, H is semi-positive definite. \square

I employed the same code I used for Problem (1); however, knowing that our function is non-convex and that H is semi-positive definite for $\forall y \geq 2$, I chose to initialize y from a normal distribution of numbers from $[2, 100]$. Although this did not guarantee that y fell below two, it handled the overflow errors that would have resulted for a different initial range.

```
import numpy as np
from numpy import linalg as LA
# Function to optimize
def function(x, y, z):
    return (-x+1)**2 + (3*y-6)**3 + (4*z+12)**2

# Line search function to find optimal learning rate
def line_search(x, y, z, grad_x, grad_y, grad_z, lr, function, alpha, beta,
max_iterations=1000, threshold=1e-12):
    for i in range(max_iterations):
        cost = function(x - lr*grad_x, y - lr*grad_y, z - lr*grad_z)
        prev_cost = function(x, y, z)
        if cost <= prev_cost - threshold*lr*LA.norm([grad_x, grad_y, grad_z]):
            return lr * beta
```

```

        else:
            return lr
    return lr

# Gradient descent with line search
def gradient_descent(x0, y0, z0, lr, iterations, threshold, function):
    history_x = [x0]
    history_y = [y0]
    history_z = [z0]
    x = x0
    y = y0
    z = z0
    step_count = 0
    alpha = .01
    beta = .05
    for i in range(iterations):
        grad_x = np.float128(-2*(x-1))
        grad_y = np.float128(9*(y-2)**2)
        grad_z = np.float128(8*(z+3))
        lr = line_search(x, y, z, grad_x, grad_y, grad_z, lr, function, alpha, beta)
        x = x - lr*grad_x
        y = y - lr*grad_y
        z = z - lr*grad_z
        history_x.append(x)
        history_y.append(y)
        history_z.append(z)
        step_count += 1
        if lr * LA.norm([grad_x, grad_y, grad_z]) < threshold: # stopping condition
            break
    print("Number of steps taken:", step_count)
    print("Minima found at:(", x,',',y,',',z,') with f(x,y,z)=',function(x,y,z))
    return history_x, history_y, history_z

# Define initial parameters
x0 = np.random.uniform(0, 50) # starting point for x
y0 = np.random.uniform(5,100) # starting point for y
z0 = np.random.uniform(-50,50) # starting point for z
lr = 0.1 # initial learning rate
iterations = 1000 # number of iterations
threshold = 1e-15 # threshold for stopping condition

# Run gradient descent with line search
history_x, history_y, history_z = gradient_descent(x0, y0, z0, lr, iterations,
    threshold, function)

```

Which produced output

Number of steps taken: 16

Minima found at:(26.783149920160626152 , -613.32636881404249407 , -1.4514259468715789799)
 with f(x,y,z)= -6290429437.343050923